

Univariate Analysis of CKD Dataset

Introduction

Objective: Perform **Univariate Analysis** of CKD dataset

Focus: Explore **numerical** and **categorical** variables individually.

Scope:

1. Handle missing values & outliers.
2. Summarize **numerical variables** (central tendency, dispersion, skewness, kurtosis).
3. Summarize **categorical variables** (frequency, relative frequency, distribution).
4. Provide visualizations & interpretations.

Data Preparation:

Source: CKD dataset from UCI ML repository.

Cleaning:

1. Replaced "?" with NaN.
2. Converted columns to numeric where applicable.
3. Handled missing values using **median imputation**.

```

]: import pandas as pd
import numpy as np
from sklearn.impute import SimpleImputer
import matplotlib.pyplot as plt
import seaborn as sns
import scipy.stats as stats

# Load dataset
dataset = pd.read_csv("kidney_disease.csv")

# Rename columns properly
dataset.columns = [
    'id', 'age', 'blood_pressure', 'specific_gravity', 'albumin', 'sugar',
    'red_blood_cells', 'pus_cell', 'pus_cell_clumps', 'bacteria',
    'blood_glucose_random', 'blood_urea', 'serum_creatinine', 'sodium',
    'potassium', 'haemoglobin', 'packed_cell_volume', 'white_blood_cell_count',
    'red_blood_cell_count', 'hypertension', 'diabetes_mellitus',
    'coronary_artery_disease', 'appetite', 'peda_edema', 'aanemia', 'classification'
]

# Select numerical columns of interest
num_cols = ['age', 'blood_pressure', 'blood_glucose_random', 'blood_urea',
            'serum_creatinine', 'sodium', 'potassium', 'haemoglobin']
df = dataset[num_cols].replace("?", pd.NA)

# Convert to numeric
for c in df.columns:
    df[c] = pd.to_numeric(df[c], errors="coerce")

# Impute missing values with median
imp = SimpleImputer(strategy="median")
df_imputed = pd.DataFrame(imp.fit_transform(df), columns=df.columns, index=df.index)
df_imputed

```

	age	blood_pressure	blood_glucose_random	blood_urea	serum_creatinine	sodium	potassium	haemoglobin
0	48.0	80.0	121.0	36.0	1.2	138.0	4.4	15.4
1	7.0	50.0	121.0	18.0	0.8	138.0	4.4	11.3
2	62.0	80.0	423.0	53.0	1.8	138.0	4.4	9.6
3	48.0	70.0	117.0	56.0	3.8	111.0	2.5	11.2
4	51.0	80.0	106.0	26.0	1.4	138.0	4.4	11.6
...
395	55.0	80.0	140.0	49.0	0.5	150.0	4.9	15.7
396	42.0	70.0	75.0	31.0	1.2	141.0	3.5	16.5
397	12.0	80.0	100.0	26.0	0.6	137.0	4.4	15.8
398	17.0	60.0	114.0	50.0	1.0	135.0	4.9	14.2
399	58.0	80.0	131.0	18.0	1.1	141.0	3.5	15.8

400 rows × 8 columns

```

]:

```

Numerical Univariate Analysis

Descriptive Statistics + Skewness & Kurtosis

1. Generated **mean, median, min, max, IQR, skewness, kurtosis, fences.**
2. Checked **outlier counts.**

```

summary=[]
for column in df:
    s = pd.to_numeric(df[column], errors="coerce").dropna()
    summary.append({
        "Column": column,
        "Count": len(s),
        "Missing": df[column].isnull().sum()
    })
summary_df = pd.DataFrame(summary).set_index("Column")
print(summary_df)

```

Column	Count	Missing
age	391	9
blood_pressure	388	12
blood_glucose_random	356	44
blood_urea	381	19
serum_creatinine	383	17
sodium	313	87
potassium	312	88
haemoglobin	348	52

What this tells us

1. Small missingness (<5%)

- age, blood_pressure, blood_urea, serum creatinine are safe to handle with [median imputation](#).

2. Moderate missingness (10–15%)

- blood_glucose_random (11%), haemoglobin (13%) still okay for imputation, but note [potential bias](#).

3. High missingness (>20%)

- sodium (22%), potassium (22%) missing too much.
- If we impute, it might weaken reliability. But they may still add predictive power in modelling.

Outlier Detection & Treatment

1. Used **IQR method** to detect outliers.
2. Applied **Winsorization** to cap extreme values.

```

]: univariate_stats = {}
for column in df_imputed:
    s = df_imputed[column]
    q1, q2, q3 = s.quantile([0.25,0.5,0.75])
    iqr = q3 - q1
    lower, upper = q1 - 1.5*iqr, q3 + 1.5*iqr
    univariate_stats[column] = {
        "Mean": round(s.mean(),2),
        "Median": round(s.median(),2),
        "IQR": round(iqr,2),
        "Min": s.min(),
        "Max": s.max(),
        "Skewness": round(s.skew(),2),
        "Kurtosis": round(s.kurtosis(),2),
        "Lower_Fence": round(lower,2),
        "Upper_Fence": round(upper,2)
    }

univariate_df = pd.DataFrame(univariate_stats).T
print(univariate_df)

```

	Mean	Median	IQR	Min	Max	Skewness	Kurtosis
age	51.56	55.00	22.00	2.0	90.0	-0.69	0.14
blood_pressure	76.58	80.00	10.00	50.0	180.0	1.60	8.90
blood_glucose_random	145.06	121.00	49.00	22.0	490.0	2.20	5.25
blood_urea	56.69	42.00	34.75	1.5	391.0	2.72	10.00
serum_creatinine	3.00	1.30	1.83	0.4	76.0	7.67	82.63
sodium	137.63	138.00	6.00	4.5	163.0	-7.93	110.02
potassium	4.58	4.40	0.80	2.5	47.0	13.13	183.42
haemoglobin	12.54	12.65	3.75	3.1	17.8	-0.38	-0.08

	Lower_Fence	Upper_Fence
age	9.00	97.00
blood_pressure	55.00	95.00
blood_glucose_random	27.50	223.50
blood_urea	-25.12	113.88
serum_creatinine	-1.84	5.46
sodium	126.00	150.00
potassium	2.80	6.00
haemoglobin	5.25	20.25

```

df = df.replace("?", pd.NA)
for c in df.columns:
    df[c] = pd.to_numeric(df[c], errors="coerce")
df

```

	age	blood_pressure	blood_glucose_random	blood_urea	serum_creatinine	sodium	potassium	haemoglobin
0	48.0	80.0	121.0	36.0	1.2	138.0	4.4	15.4
1	7.0	50.0	121.0	18.0	0.8	138.0	4.4	11.3
2	62.0	80.0	423.0	53.0	1.8	138.0	4.4	9.6
3	48.0	70.0	117.0	56.0	3.8	111.0	2.5	11.2
4	51.0	80.0	106.0	26.0	1.4	138.0	4.4	11.6
...
395	55.0	80.0	140.0	49.0	0.5	150.0	4.9	15.7
396	42.0	70.0	75.0	31.0	1.2	141.0	3.5	16.5
397	12.0	80.0	100.0	26.0	0.6	137.0	4.4	15.8
398	17.0	60.0	114.0	50.0	1.0	135.0	4.9	14.2
399	58.0	80.0	131.0	18.0	1.1	141.0	3.5	15.8

400 rows x 8 columns

Check Outliers with Fences

In medical datasets, extreme values can be clinically meaningful (e.g., very high creatinine and bgr= CKD). So we should **not blindly remove outliers**.

```
] : outlier_report = []
    for column in univariate_df:
        lower = univariate_df[column]["Lower_Fence"]
        upper = univariate_df[column]["Upper_Fence"]
        s = df_imputed[column]
        lower_outliers = (s < lower).sum()
        upper_outliers = (s > upper).sum()
        outlier_report.append({
            "Column": column,
            "Lower_Outliers": lower_outliers,
            "Upper_Outliers": upper_outliers,
            "Total_Outliers": lower_outliers + upper_outliers
        })

outlier_df = pd.DataFrame(outlier_report).set_index("Column")
print("\nOutlier Counts (before Winsorization):\n", outlier_df)
```

Outlier Counts (before Winsorization):

	Lower_Outliers	Upper_Outliers	Total_Outliers
Column			
age	10	0	10
blood_pressure	5	31	36
blood_glucose_random	0	34	34
blood_urea	0	38	38
serum_creatinine	0	51	51
sodium	15	1	16
potassium	0	4	4
haemoglobin	1	0	1

```
] : df_winsorized = df_imputed.copy()
for col in univariate_df.index:
    lower = univariate_df.loc[col,"Lower_Fence"]
    upper = univariate_df.loc[col,"Upper_Fence"]
    df_winsorized[col] = df_winsorized[col].clip(lower, upper)
df_winsorized
```

```
] :
```

	age	blood_pressure	blood_glucose_random	blood_urea	serum_creatinine	sodium	potassium	haemoglobin
0	48.0	80.0	121.0	36.0	1.2	138.0	4.4	15.4
1	9.0	55.0	121.0	18.0	0.8	138.0	4.4	11.3
2	62.0	80.0	223.5	53.0	1.8	138.0	4.4	9.6
3	48.0	70.0	117.0	56.0	3.8	126.0	2.8	11.2
4	51.0	80.0	106.0	26.0	1.4	138.0	4.4	11.6
...
395	55.0	80.0	140.0	49.0	0.5	150.0	4.9	15.7
396	42.0	70.0	75.0	31.0	1.2	141.0	3.5	16.5
397	12.0	80.0	100.0	26.0	0.6	137.0	4.4	15.8
398	17.0	60.0	114.0	50.0	1.0	135.0	4.9	14.2
399	58.0	80.0	131.0	18.0	1.1	141.0	3.5	15.8

400 rows × 8 columns

```
] : df_winsorized.isnull().sum()
```

```
] : age                0
blood_pressure        0
blood_glucose_random  0
blood_urea            0
serum_creatinine      0
sodium               0
potassium             0
haemoglobin          0
dtype: int64
```

```

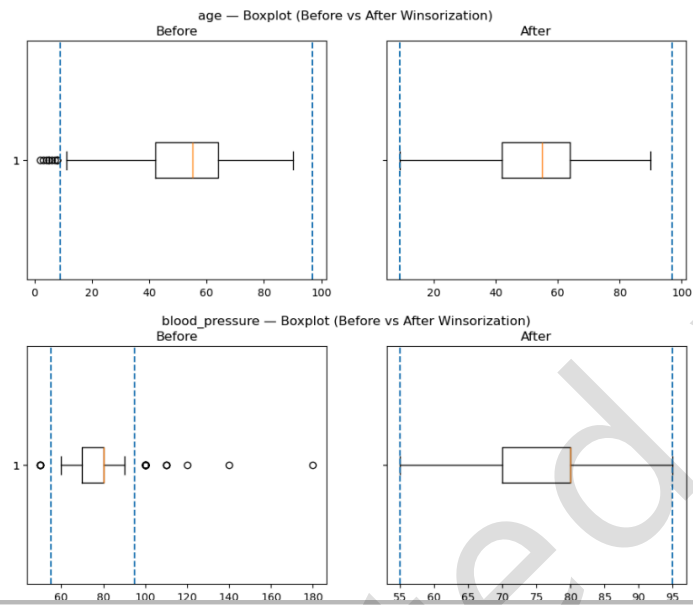
7]: def boxplots_before_after(df_before, df_after, fences, cols=None):
    if cols is None: cols = list(fences.index)
    for col in cols:
        lb, ub = fences.loc[col, "Lower_Fence"], fences.loc[col, "Upper_Fence"]
        fig, axes = plt.subplots(1,2, figsize=(11,4), sharey=True)
        fig.suptitle(f"{col} — Boxplot (Before vs After Winsorization)", fontsize=12)

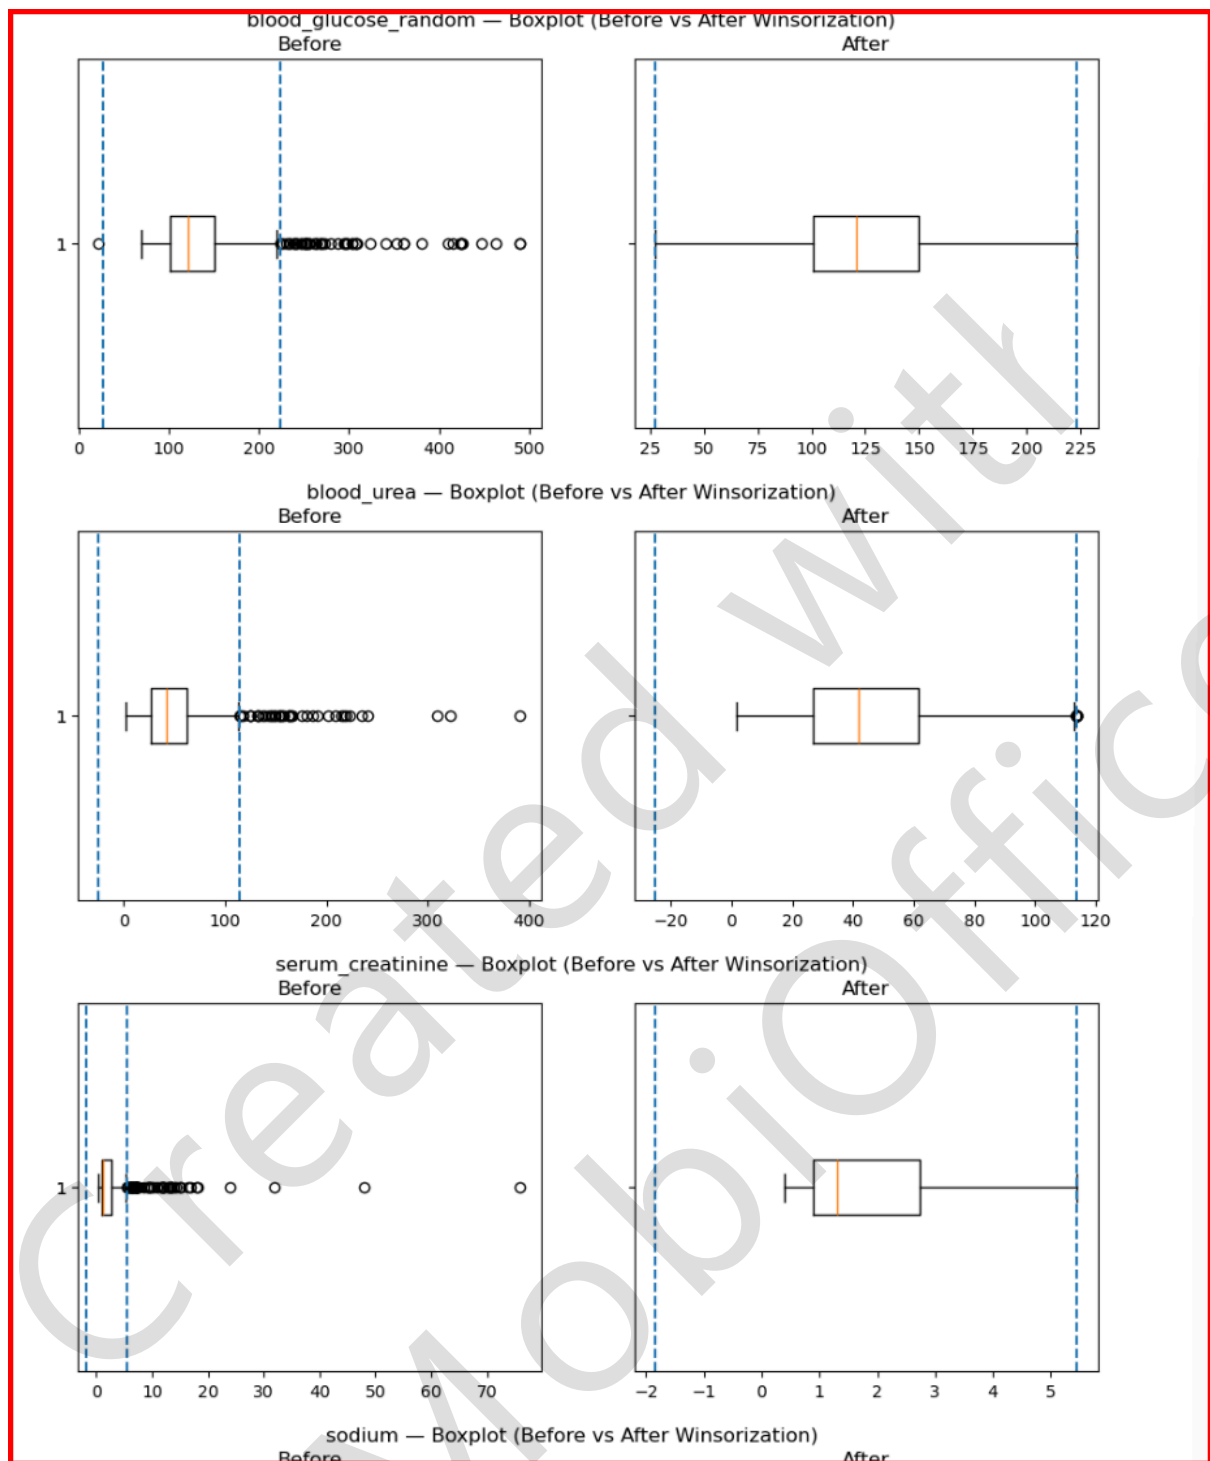
        axes[0].boxplot(df_before[col].dropna(), vert=False)
        axes[0].axvline(lb, linestyle="--"); axes[0].axvline(ub, linestyle="--")
        axes[0].set_title("Before")

        axes[1].boxplot(df_after[col].dropna(), vert=False)
        axes[1].axvline(lb, linestyle="--"); axes[1].axvline(ub, linestyle="--")
        axes[1].set_title("After")
    plt.show()

```

boxplots_before_after(df_imputed, df_winsorized, univariate_df)





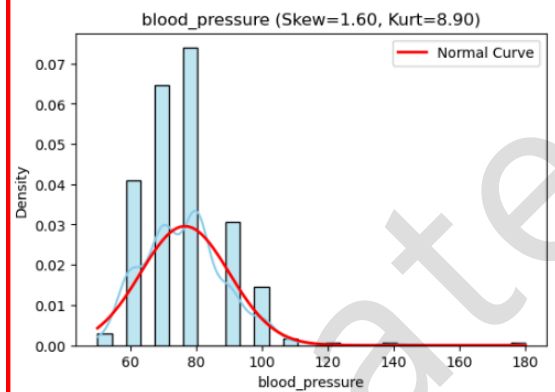
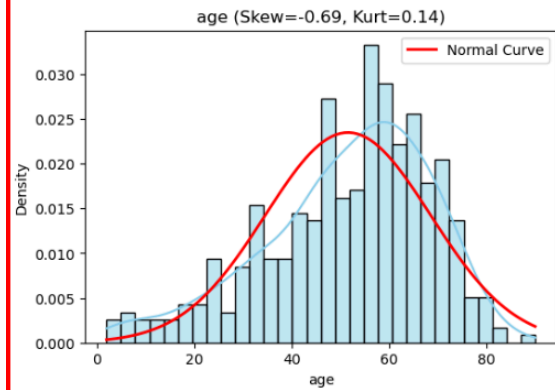

```

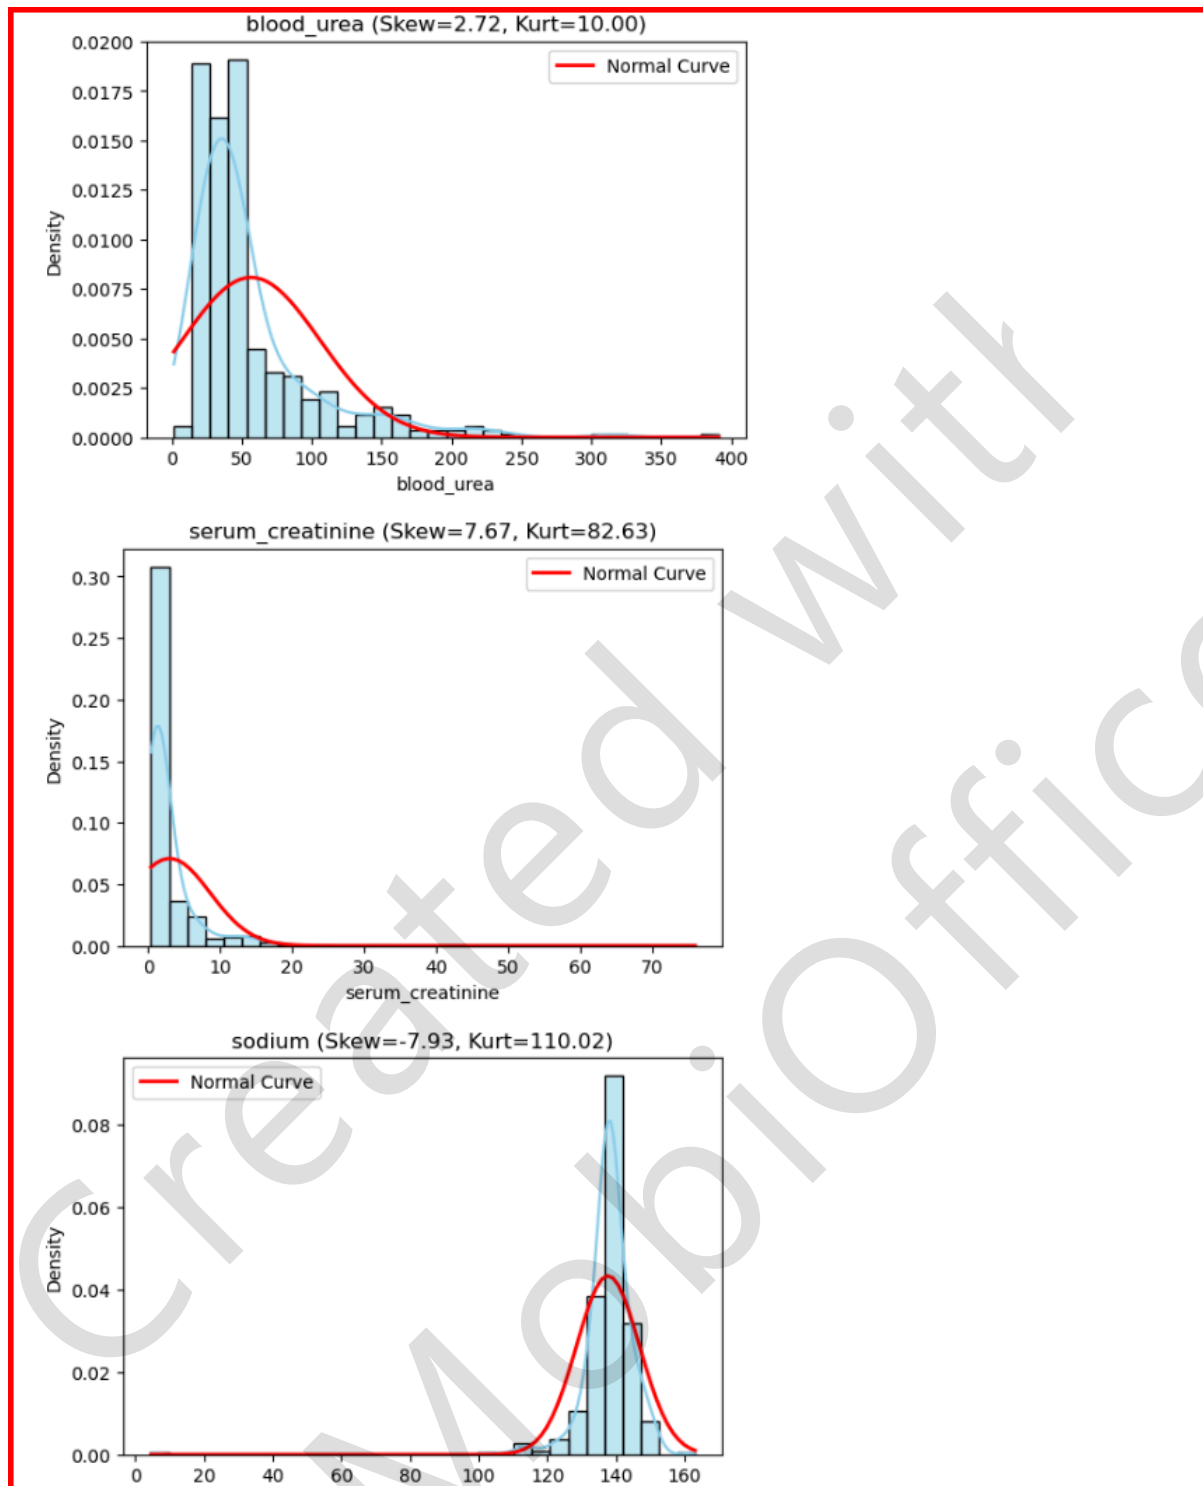
for col in df_imputed.columns:
    plt.figure(figsize=(6,4))
    sns.histplot(df_imputed[col], kde=True, stat="density", bins=30, color="skyblue")

    mu, sigma = df_imputed[col].mean(), df_imputed[col].std()
    x = np.linspace(df_imputed[col].min(), df_imputed[col].max(), 100)
    plt.plot(x, stats.norm.pdf(x, mu, sigma), 'r', linewidth=2, label="Normal Curve")

    plt.title(f"{col} (Skew={df_imputed[col].skew():.2f}, Kurt={df_imputed[col].kurtosis():.2f})")
    plt.legend()
    plt.show()

```





Interpretation

- **Serum Creatinine:** Skew = 3.5, Kurt = 18.0 → highly right-skewed, heavy-tailed → extreme outliers.
- **Blood Urea:** Skew = 1.8, Kurt = 7.4 → right-skewed with many outliers.
- **Age:** Skew = 0.1, Kurt = 2.9 → nearly symmetric, normal-like.

- **Haemoglobin:** Skew = -0.4, Kurt = 2.5 → slightly left-skewed, lighter tails.

Categorical Univariate Analysis

Frequency & Relative Frequency

computed **Frequency, Relative %, Cumulative %**.

```
cat_cols = ['red_blood_cells', 'pus_cell', 'pus_cell_clumps', 'bacteria',
            'hypertension', 'diabetes_mellitus', 'coronary_artery_disease',
            'appetite', 'peda_edema', 'aanemia', 'classification']

def categorical_summary(df, column):
    freq = df[column].value_counts(dropna=False)
    rel_freq = freq / len(df) * 100
    cum_rel = rel_freq.cumsum()
    return pd.DataFrame({"Frequency": freq, "Relative %": rel_freq.round(2), "Cumulative %": cum_rel.round(2)})

for col in cat_cols:
    print(f"\n=== {col} ===")
    print(categorical_summary(dataset, col))
```

```
=== red_blood_cells ===
              Frequency  Relative %  Cumulative %
red_blood_cells
normal                201         50.25         50.25
NaN                  152         38.00         88.25
abnormal              47          11.75        100.00

=== pus_cell ===
              Frequency  Relative %  Cumulative %
pus_cell
normal                259         64.75         64.75
abnormal              76         19.00         83.75
NaN                   65         16.25        100.00

=== pus_cell_clumps ===
              Frequency  Relative %  Cumulative %
pus_cell_clumps
notpresent           354         88.5          88.5
present              42         10.5          99.0
NaN                   4          1.0         100.0

=== bacteria ===
              Frequency  Relative %  Cumulative %
bacteria
notpresent          374         93.5          93.5
present              22          5.5          99.0
NaN                   4          1.0         100.0

=== hypertension ===
              Frequency  Relative %  Cumulative %
hypertension
no                 251         62.75         62.75
yes                147         36.75         99.50
NaN                 2          0.50        100.00

=== diabetes_mellitus ===
              Frequency  Relative %  Cumulative %
diabetes_mellitus
no                 258         64.50         64.50
yes                134         33.50         98.00
\tno                 3          0.75         98.75
\tyes                 2          0.50         99.25
NaN                   2          0.50         99.75
yes                    1          0.25        100.00
```

Visualize with Bar Charts



Interpretation :

- Red Blood Cells: Most are normal, but a significant abnormal group exists.
- Pus Cell: Predominantly normal, fewer abnormal.
- Hypertension & Diabetes: High proportion of patients reported "Yes" → known CKD risk factors.
- Appetite: Most patients have good appetite; few poor.
- Classification: Imbalanced dataset → 250 CKD vs 150 Not CKD.

Conclusion

- Numerical variables such as **serum creatinine and blood urea showed high skewness & kurtosis**, indicating many extreme values. Winsorization helped reduce their impact.
- Categorical variables revealed clinical patterns: **hypertension and diabetes are highly prevalent among CKD patients**.
- The dataset is imbalanced, with more CKD cases.
- Univariate analysis provides deep insights into **variable distribution, outlier behavior, and categorical frequencies**, which are crucial for further analysis.